

Course 2

Intermediate PyTorch

Comprehensive Course Notes

Module 1: Hyperparameter Optimization with Optuna

Module 2: Computer Vision & Transfer Learning

Module 3: NLP with Transformers

Module 4: PyTorch Lightning

Contents

1	Module 1: Hyperparameter Optimization with Optuna	2
1.1	Introduction to AutoML	2
1.2	Optuna Framework	2
1.3	Suggest Methods	3
1.4	FlexibleCNN Architecture	3
1.5	Pruning Unpromising Trials	4
1.6	Study Persistence	4
1.7	Visualization	5
2	Module 2: Computer Vision & Transfer Learning	6
2.1	What is Transfer Learning?	6
2.2	TorchVision Models	6
2.3	Modifying for Custom Tasks	6
2.4	Freezing Pretrained Weights	7
2.5	ImageFolder Dataset	7
2.6	ImageNet Preprocessing	8
2.7	Learning Rate Scheduling	8
2.8	Differential Learning Rates	9
3	Module 3: NLP with Transformers	10
3.1	Text Classification Pipeline	10
3.2	Tokenization	10
3.3	DistilBERT for Classification	10
3.4	Custom Text Classifier	11
3.5	Custom Dataset for Text	11
3.6	Handling Imbalanced Classes	12
3.7	Fine-Tuning Strategies	12
3.8	Word Embeddings Concepts	12
4	Module 4: PyTorch Lightning	13
4.1	Why PyTorch Lightning?	13
4.2	LightningModule Structure	13
4.3	LightningDataModule	14
4.4	Trainer Configuration	14
4.5	Callbacks	15
4.6	Learning Rate Scheduling	15
4.7	Logging Metrics	16
4.8	Loading Checkpoints	16
5	Quiz Review: Key Concepts	18
5.1	Quiz 5: Hyperparameter Optimization	18
5.2	Quiz 6: Computer Vision	18
5.3	Quiz 7: NLP Fundamentals	18
5.4	Quiz 8: PyTorch Lightning	19
6	Quick Reference	20
6.1	Optuna Template	20
6.2	Transfer Learning Template	20
6.3	Lightning Template	20

1 Module 1: Hyperparameter Optimization with Optuna

1.1 Introduction to AutoML

What is AutoML?

AutoML (Automated Machine Learning) automates the machine learning pipeline:

- Feature engineering
- Model selection
- **Hyperparameter optimization** ← Focus of this module
- Architecture search

1.2 Optuna Framework

Optuna Key Concepts

- **Study:** A complete optimization session
- **Trial:** A single training run with specific hyperparameters
- **Objective Function:** Returns the metric to optimize
- **Sampler:** Strategy for selecting hyperparameters (TPE, Random, etc.)
- **Pruner:** Early stops unpromising trials

Listing 1: Basic Optuna Setup

```
import optuna

def objective(trial):
    # Suggest hyperparameters
    lr = trial.suggest_float("lr", 1e-5, 1e-1, log=True)
    n_layers = trial.suggest_int("n_layers", 1, 5)
    dropout = trial.suggest_float("dropout", 0.1, 0.5)

    # Build model with suggested params
    model = build_model(n_layers, dropout)

    # Train and evaluate
    accuracy = train_and_evaluate(model, lr)

    return accuracy # Optuna will maximize/minimize this

# Create and run study
study = optuna.create_study(direction="maximize")
study.optimize(objective, n_trials=100)

# Best hyperparameters
print(study.best_params)
print(study.best_value)
```

1.3 Suggest Methods

Method	Usage
<code>suggest_int(name, low, high)</code>	Integer values
<code>suggest_float(name, low, high)</code>	Float values
<code>suggest_float(..., log=True)</code>	Log-scale sampling (for learning rates)
<code>suggest_categorical(name, choices)</code>	Choose from list of options

Listing 2: Hyperparameter Suggestions

```
def objective(trial):
    # Learning rate (log scale)
    lr = trial.suggest_float("lr", 1e-5, 1e-1, log=True)

    # Architecture choices
    n_filters = trial.suggest_int("n_filters", 16, 128)
    n_layers = trial.suggest_int("n_layers", 2, 6)

    # Regularization
    dropout = trial.suggest_float("dropout", 0.0, 0.5)

    # Categorical choice
    optimizer_name = trial.suggest_categorical(
        "optimizer", ["Adam", "SGD", "RMSprop"])
)

# Build and train model...
```

1.4 FlexibleCNN Architecture

Listing 3: Flexible CNN for Hyperparameter Search

```
class FlexibleCNN(nn.Module):
    def __init__(self, n_layers, n_filters, dropout_rate, num_classes):
        super().__init__()

        layers = []
        in_channels = 3

        for i in range(n_layers):
            out_channels = n_filters * (2 ** i)  # Double channels
            layers.extend([
                nn.Conv2d(in_channels, out_channels, 3, padding=1),
                nn.BatchNorm2d(out_channels),
                nn.ReLU(),
                nn.MaxPool2d(2, 2)
            ])
            in_channels = out_channels

        self.features = nn.Sequential(*layers)
        self.flatten = nn.Flatten()

        # Calculate flattened size
        with torch.no_grad():
            dummy = torch.zeros(1, 3, 32, 32)
            flat_size = self.features(dummy).view(1, -1).size(1)
```

```

        self.classifier = nn.Sequential(
            nn.Linear(flat_size, 256),
            nn.ReLU(),
            nn.Dropout(dropout_rate),
            nn.Linear(256, num_classes)
        )

    def forward(self, x):
        x = self.features(x)
        x = self.flatten(x)
        return self.classifier(x)

```

1.5 Pruning Unpromising Trials

Median Pruner

Stops trials that perform worse than the median of completed trials at the same step:

```

study = optuna.create_study(
    direction="maximize",
    pruner=optuna.pruners.MedianPruner(
        n_startup_trials=5,
        n_warmup_steps=5
    )
)

```

Listing 4: Reporting for Pruning

```

def objective(trial):
    model = build_model(trial)

    for epoch in range(num_epochs):
        train_loss = train_one_epoch(model)
        val_acc = evaluate(model)

        # Report intermediate value for pruning
        trial.report(val_acc, epoch)

        # Check if trial should be pruned
        if trial.should_prune():
            raise optuna.TrialPruned()

    return val_acc

```

1.6 Study Persistence

Listing 5: Saving and Loading Studies

```

# Save to SQLite database
study = optuna.create_study(
    study_name="cnn_optimization",
    storage="sqlite:///optuna_study.db",
    load_if_exists=True # Resume existing study
)

# Continue optimization
study.optimize(objective, n_trials=50)

```

```
# Analyze results
print(f"Best trial: {study.best_trial.number}")
print(f"Best params: {study.best_params}")
print(f"Best value: {study.best_value}")

# Get all trials
df = study.trials_dataframe()
```

1.7 Visualization

Listing 6: Optuna Visualization

```
import optuna.visualization as vis

# Parameter importance
vis.plot_param_importances(study)

# Optimization history
vis.plot_optimization_history(study)

# Parameter relationships
vis.plot_parallel_coordinate(study)

# Hyperparameter slice
vis.plot_slice(study)
```

2 Module 2: Computer Vision & Transfer Learning

2.1 What is Transfer Learning?

Transfer Learning

Using a model pretrained on a large dataset (e.g., ImageNet) and adapting it to a new task:

- Leverages learned features (edges, textures, objects)
- Requires less data for the new task
- Faster training convergence
- Often achieves better results than training from scratch

2.2 TorchVision Models

Listing 7: Loading Pretrained Models

```
import torchvision.models as models

# Load pretrained model
model = models.mobilenet_v3_small(weights='IMAGENET1K_V1')

# Alternative models
resnet18 = models.resnet18(weights='IMAGENET1K_V1')
efficientnet = models.efficientnet_b0(weights='IMAGENET1K_V1')
vgg16 = models.vgg16(weights='IMAGENET1K_V1')
```

Model Selection Considerations

- **MobileNetV3:** Lightweight, fast, mobile-friendly
- **ResNet:** Reliable, well-understood, various depths
- **EfficientNet:** Best accuracy/efficiency tradeoff
- **VGG:** Simple architecture, good for visualization

2.3 Modifying for Custom Tasks

Listing 8: Replacing the Classifier

```
# MobileNetV3
model = models.mobilenet_v3_small(weights='IMAGENET1K_V1')
num_classes = 10

# Get number of input features to classifier
in_features = model.classifier[0].in_features

# Replace classifier
model.classifier = nn.Sequential(
    nn.Linear(in_features, 256),
    nn.Hardswish(),
    nn.Dropout(p=0.2),
    nn.Linear(256, num_classes)
)
```

```
# ResNet (replaces fc layer)
resnet = models.resnet18(weights='IMAGENET1K_V1')
resnet.fc = nn.Linear(resnet.fc.in_features, num_classes)
```

2.4 Freezing Pretrained Weights

Listing 9: Feature Extraction vs Fine-Tuning

```
# Option 1: Feature Extraction (freeze all)
for param in model.parameters():
    param.requires_grad = False

# Unfreeze classifier only
for param in model.classifier.parameters():
    param.requires_grad = True

# Option 2: Gradual Unfreezing (fine-tune later layers)
# Freeze early layers
for name, param in model.named_parameters():
    if 'features.0' in name or 'features.1' in name:
        param.requires_grad = False
```

Freezing Strategy

1. Start with frozen backbone, train classifier only
2. Gradually unfreeze later layers
3. Use lower learning rate for pretrained weights

2.5 ImageFolder Dataset

Listing 10: Using ImageFolder

```
from torchvision.datasets import ImageFolder

# Directory structure:
# data/
#   train/
#     class_a/
#     class_b/
#   val/
#     class_a/
#     class_b/

train_dataset = ImageFolder(
    root='data/train',
    transform=train_transform
)

val_dataset = ImageFolder(
    root='data/val',
    transform=val_transform
)

# Class names
```

```
print(train_dataset.classes) # ['class_a', 'class_b']
print(train_dataset.class_to_idx) # {'class_a': 0, 'class_b': 1}
```

2.6 ImageNet Preprocessing

Listing 11: Standard Preprocessing for Pretrained Models

```
# ImageNet statistics
IMAGENET_MEAN = [0.485, 0.456, 0.406]
IMAGENET_STD = [0.229, 0.224, 0.225]

train_transform = transforms.Compose([
    transforms.Resize(256),
    transforms.RandomResizedCrop(224),
    transforms.RandomHorizontalFlip(),
    transforms.ToTensor(),
    transforms.Normalize(IMAGENET_MEAN, IMAGENET_STD)
])

val_transform = transforms.Compose([
    transforms.Resize(256),
    transforms.CenterCrop(224),
    transforms.ToTensor(),
    transforms.Normalize(IMAGENET_MEAN, IMAGENET_STD)
])
```

2.7 Learning Rate Scheduling

Listing 12: Learning Rate Schedulers

```
from torch.optim.lr_scheduler import StepLR, ReduceLROnPlateau

# Step decay
scheduler = StepLR(optimizer, step_size=10, gamma=0.1)
# Reduce LR by 0.1 every 10 epochs

# Adaptive (reduce on plateau)
scheduler = ReduceLROnPlateau(
    optimizer,
    mode='min',          # Minimize metric
    factor=0.1,           # Multiply LR by this
    patience=5,            # Wait this many epochs
    verbose=True
)

# In training loop
for epoch in range(num_epochs):
    train_loss = train_epoch(...)
    val_loss = validate(...)

    # For StepLR
    scheduler.step()

    # For ReduceLROnPlateau
    scheduler.step(val_loss)
```

2.8 Differential Learning Rates

Listing 13: Different LR for Different Layers

```
# Lower LR for pretrained features, higher for new classifier
optimizer = optim.Adam([
    {'params': model.features.parameters(), 'lr': 1e-5},
    {'params': model.classifier.parameters(), 'lr': 1e-3}
])
```

3 Module 3: NLP with Transformers

3.1 Text Classification Pipeline

NLP Pipeline Steps

1. **Tokenization:** Convert text to tokens (words/subwords)
2. **Encoding:** Convert tokens to numerical IDs
3. **Embedding:** Map IDs to dense vectors
4. **Processing:** Apply transformer/RNN layers
5. **Classification:** Final prediction layer

3.2 Tokenization

Listing 14: Basic Tokenization

```
from transformers import AutoTokenizer

tokenizer = AutoTokenizer.from_pretrained('distilbert-base-uncased')

# Tokenize text
text = "I love PyTorch!"
tokens = tokenizer(
    text,
    padding='max_length',
    truncation=True,
    max_length=128,
    return_tensors='pt'
)

print(tokens['input_ids'].shape)      # [1, 128]
print(tokens['attention_mask'].shape) # [1, 128]
```

Tokenizer Output

- **input_ids:** Token indices for the vocabulary
- **attention_mask:** 1 for real tokens, 0 for padding
- **Special tokens:** [CLS] at start, [SEP] at end

3.3 DistilBERT for Classification

Listing 15: Loading DistilBERT

```
from transformers import DistilBertModel, DistilBertConfig

# Load pretrained DistilBERT
bert = DistilBertModel.from_pretrained('distilbert-base-uncased')

# Get hidden size
hidden_size = bert.config.hidden_size # 768
```

3.4 Custom Text Classifier

Listing 16: Text Classification Model

```

class TextClassifier(nn.Module):
    def __init__(self, num_classes, dropout=0.3):
        super().__init__()

        self.bert = DistilBertModel.from_pretrained(
            'distilbert-base-uncased'
        )
        self.dropout = nn.Dropout(dropout)
        self.classifier = nn.Linear(
            self.bert.config.hidden_size, # 768
            num_classes
        )

    def forward(self, input_ids, attention_mask):
        # Get BERT output
        outputs = self.bert(
            input_ids=input_ids,
            attention_mask=attention_mask
        )

        # Use [CLS] token representation
        cls_output = outputs.last_hidden_state[:, 0, :]

        # Classify
        x = self.dropout(cls_output)
        logits = self.classifier(x)
        return logits

```

3.5 Custom Dataset for Text

Listing 17: Text Dataset Class

```

class TextDataset(Dataset):
    def __init__(self, texts, labels, tokenizer, max_length=128):
        self.texts = texts
        self.labels = labels
        self.tokenizer = tokenizer
        self.max_length = max_length

    def __len__(self):
        return len(self.texts)

    def __getitem__(self, idx):
        encoding = self.tokenizer(
            self.texts[idx],
            padding='max_length',
            truncation=True,
            max_length=self.max_length,
            return_tensors='pt'
        )

        return {
            'input_ids': encoding['input_ids'].squeeze(0),
            'attention_mask': encoding['attention_mask'].squeeze(0),
            'label': torch.tensor(self.labels[idx], dtype=torch.long)
        }

```

}

3.6 Handling Imbalanced Classes

Listing 18: Weighted Cross-Entropy

```
# Calculate class weights (inverse frequency)
class_counts = [1000, 200, 50] # Example counts
weights = [1.0 / c for c in class_counts]
weights = torch.tensor(weights) / sum(weights) * len(weights)

# Weighted loss
criterion = nn.CrossEntropyLoss(weight=weights.to(device))
```

3.7 Fine-Tuning Strategies

Transformer Fine-Tuning Tips

- Use small learning rate (2e-5 to 5e-5)
- Train for few epochs (2-4 typically sufficient)
- Watch for overfitting on small datasets
- Consider freezing early layers

Listing 19: Fine-Tuning Setup

```
# Freeze BERT, train classifier only
for param in model.bert.parameters():
    param.requires_grad = False

# Or use different learning rates
optimizer = optim.AdamW([
    {'params': model.bert.parameters(), 'lr': 2e-5},
    {'params': model.classifier.parameters(), 'lr': 1e-3}
])
```

3.8 Word Embeddings Concepts

Embedding Layer

Maps discrete tokens to continuous vectors:

$$\text{embedding}(x) = W_e[x]$$

Where W_e is the embedding matrix of shape (vocab_size, embedding_dim)

```
embedding = nn.Embedding(
    num_embeddings=30000, # Vocabulary size
    embedding_dim=768 # Vector dimension
)
```

4 Module 4: PyTorch Lightning

4.1 Why PyTorch Lightning?

Lightning Benefits

- Removes boilerplate training code
- Built-in best practices (gradient clipping, logging, checkpointing)
- Easy multi-GPU training
- Clean, organized code structure
- Focus on research, not engineering

4.2 LightningModule Structure

Listing 20: Basic LightningModule

```
import pytorch_lightning as pl

class LitModel(pl.LightningModule):
    def __init__(self, num_classes, learning_rate=1e-3):
        super().__init__()
        self.save_hyperparameters() # Log all __init__ args

        # Define model
        self.model = nn.Sequential(
            nn.Linear(784, 256),
            nn.ReLU(),
            nn.Linear(256, num_classes)
        )
        self.criterion = nn.CrossEntropyLoss()

    def forward(self, x):
        return self.model(x)

    def training_step(self, batch, batch_idx):
        x, y = batch
        logits = self(x)
        loss = self.criterion(logits, y)

        # Logging (automatic tensorboard)
        self.log('train_loss', loss, prog_bar=True)
        return loss

    def validation_step(self, batch, batch_idx):
        x, y = batch
        logits = self(x)
        loss = self.criterion(logits, y)

        # Calculate accuracy
        preds = logits.argmax(dim=1)
        acc = (preds == y).float().mean()

        self.log('val_loss', loss, prog_bar=True)
        self.log('val_acc', acc, prog_bar=True)
```

```

def configure_optimizers(self):
    optimizer = torch.optim.Adam(
        self.parameters(),
        lr=self.hparams.learning_rate
    )
    return optimizer

```

4.3 LightningDataModule

Listing 21: Data Module Structure

```

class MNISTDataModule(pl.LightningDataModule):
    def __init__(self, batch_size=64, data_dir='./data'):
        super().__init__()
        self.batch_size = batch_size
        self.data_dir = data_dir
        self.transform = transforms.Compose([
            transforms.ToTensor(),
            transforms.Normalize((0.1307,), (0.3081,))
        ])

    def prepare_data(self):
        # Download only (called once on main process)
        MNIST(self.data_dir, train=True, download=True)
        MNIST(self.data_dir, train=False, download=True)

    def setup(self, stage=None):
        # Called on every process
        if stage == 'fit' or stage is None:
            mnist_full = MNIST(
                self.data_dir, train=True, transform=self.transform
            )
            self.train_data, self.val_data = random_split(
                mnist_full, [55000, 5000]
            )

        if stage == 'test' or stage is None:
            self.test_data = MNIST(
                self.data_dir, train=False, transform=self.transform
            )

    def train_dataloader(self):
        return DataLoader(self.train_data, batch_size=self.batch_size,
                         shuffle=True, num_workers=4)

    def val_dataloader(self):
        return DataLoader(self.val_data, batch_size=self.batch_size,
                         num_workers=4)

    def test_dataloader(self):
        return DataLoader(self.test_data, batch_size=self.batch_size,
                         num_workers=4)

```

4.4 Trainer Configuration

Listing 22: Training with Lightning

```

from pytorch_lightning import Trainer

model = LitModel(num_classes=10)
data = MNISTDataModule(batch_size=64)

trainer = Trainer(
    max_epochs=10,
    accelerator='auto',           # Auto-detect GPU/CPU
    devices='auto',               # Use all available devices
    precision=16,                 # Mixed precision (faster)
    gradient_clip_val=1.0,        # Gradient clipping
    log_every_n_steps=50
)

# Train
trainer.fit(model, data)

# Test
trainer.test(model, data)

```

4.5 Callbacks

Listing 23: Built-in Callbacks

```

from pytorch_lightning.callbacks import (
    ModelCheckpoint, EarlyStopping, LearningRateMonitor
)

# Save best model
checkpoint_callback = ModelCheckpoint(
    monitor='val_loss',
    mode='min',
    save_top_k=3,
    filename='model-{epoch:02d}-{val_loss:.2f}',
)

# Stop if no improvement
early_stop = EarlyStopping(
    monitor='val_loss',
    patience=5,
    mode='min'
)

# Log learning rate
lr_monitor = LearningRateMonitor(logging_interval='step')

trainer = Trainer(
    callbacks=[checkpoint_callback, early_stop, lr_monitor],
    max_epochs=100
)

```

4.6 Learning Rate Scheduling

Listing 24: Adding Scheduler to LightningModule

```

def configure_optimizers(self):
    optimizer = torch.optim.Adam(self.parameters(), lr=1e-3)

```

```

        scheduler = torch.optim.lr_scheduler.ReduceLROnPlateau(
            optimizer,
            mode='min',
            factor=0.1,
            patience=3
        )

    return {
        'optimizer': optimizer,
        'lr_scheduler': {
            'scheduler': scheduler,
            'monitor': 'val_loss', # Metric to monitor
            'interval': 'epoch',
            'frequency': 1
        }
    }
}

```

4.7 Logging Metrics

Listing 25: Advanced Logging

```

def training_step(self, batch, batch_idx):
    x, y = batch
    logits = self(x)
    loss = self.criterion(logits, y)

    # Multiple metrics
    preds = logits.argmax(dim=1)
    acc = (preds == y).float().mean()

    # Log with different settings
    self.log('train_loss', loss, on_step=True, on_epoch=True)
    self.log('train_acc', acc, on_step=False, on_epoch=True)

    # Log dictionary
    self.log_dict({
        'train/loss': loss,
        'train/acc': acc
    }, prog_bar=True)

    return loss

```

4.8 Loading Checkpoints

Listing 26: Checkpoint Management

```

# Load from checkpoint
model = LitModel.load_from_checkpoint(
    'checkpoints/best-model.ckpt',
    num_classes=10
)

# Resume training
trainer = Trainer(max_epochs=20)
trainer.fit(model, data, ckpt_path='checkpoints/last.ckpt')

```

```
# Inference
model.eval()
with torch.no_grad():
    predictions = model(test_input)
```

5 Quiz Review: Key Concepts

5.1 Quiz 5: Hyperparameter Optimization

1. **AutoML scope:** Hyperparameters, architecture, feature selection
2. `suggest_float(..., log=True)` for learning rates
3. Trial pruning stops **unpromising configurations early**
4. `trial.report()` used to report intermediate results
5. Flexible CNN layer structure enables architecture search
6. `n_startup_trials` sets trials before pruner activates
7. Study results accessed via `study.best_params`
8. SQLite storage enables study persistence/resumption
9. `suggest_categorical()` for discrete choices like optimizer
10. Optuna is not neural-network specific (works with any ML)

5.2 Quiz 6: Computer Vision

1. **Transfer learning:** Uses pretrained models as starting point
2. ImageNet models trained on **1000 classes**
3. **Freezing:** Setting `requires_grad=False`
4. **ImageFolder** expects class subfolders
5. ImageNet input size: **224×224**
6. Standard normalization: `mean=[0.485, 0.456, 0.406]`
7. **Feature extraction:** Freeze backbone, train new head
8. MobileNetV3 optimized for **mobile/edge devices**
9. Learning rate schedulers adjust LR during training
10. Model surgery: freeze weights + replace classifier

5.3 Quiz 7: NLP Fundamentals

1. **Tokenization:** First step in NLP pipeline
2. Attention mask: 1=real token, 0=padding

CLS token used for **classification**

3. **Embeddings:** Dense vector representations
4. Weighted loss handles class imbalance
5. DistilBERT is a **distilled** (compressed) BERT
6. Fine-tuning adapts pretrained models to new tasks
7. Subword tokenization handles unknown words
8. Truncation handles sequences longer than `max_length`
9. `last_hidden_state` contains contextualized embeddings

5.4 Quiz 8: PyTorch Lightning

1. Lightning reduces boilerplate, provides best practices
2. **LightningModule** encapsulates model + training logic
3. `configure_optimizers()` returns optimizer (and scheduler)
4. **LightningDataModule** handles all data loading
5. `prepare_data()` for downloads, `setup()` for splits
6. EarlyStopping prevents overfitting
7. `self.log()` for automatic metric tracking
8. Trainer handles training loop, devices, precision
9. Callbacks extend functionality (checkpoints, early stopping)
10. `save_hyperparameters()` logs all `__init__` args

6 Quick Reference

6.1 Optuna Template

```
import optuna

def objective(trial):
    # Suggest hyperparameters
    lr = trial.suggest_float("lr", 1e-5, 1e-1, log=True)
    model = build_model(trial)

    for epoch in range(num_epochs):
        val_acc = train_and_validate(model, lr)
        trial.report(val_acc, epoch)
        if trial.should_prune():
            raise optuna.TrialPruned()

    return val_acc

study = optuna.create_study(
    direction="maximize",
    pruner=optuna.pruners.MedianPruner()
)
study.optimize(objective, n_trials=100)
```

6.2 Transfer Learning Template

```
import torchvision.models as models

model = models.mobilenet_v3_small(weights='IMAGENET1K_V1')

# Freeze backbone
for param in model.features.parameters():
    param.requires_grad = False

# Replace classifier
model.classifier = nn.Sequential(
    nn.Linear(576, 256),
    nn.ReLU(),
    nn.Dropout(0.2),
    nn.Linear(256, num_classes)
)
```

6.3 Lightning Template

```
class LitModel(pl.LightningModule):
    def __init__(self, num_classes, lr=1e-3):
        super().__init__()
        self.save_hyperparameters()
        self.model = ...
        self.criterion = nn.CrossEntropyLoss()

    def forward(self, x):
        return self.model(x)

    def training_step(self, batch, batch_idx):
```

```
x, y = batch
loss = self.criterion(self(x), y)
self.log('train_loss', loss)
return loss

def configure_optimizers(self):
    return torch.optim.Adam(self.parameters(),
                           lr=self.hparams.lr)

trainer = pl.Trainer(max_epochs=10, accelerator='auto')
trainer.fit(model, datamodule)
```